

# Introduction to Machine Learning

JJ Vie

2020

# Machine Learning

- ▶ Teach machine to learn from examples

Supervised learning (TP 1) & deep learning (TP 4)

Handwritten digit recognition, face recognition, etc.

Unsupervised learning (TP 2)

Recommender systems, clustering

Reinforcement learning (TP 3)

Robotics, games

# TP 1: Supervised learning

From features  $\mathbf{x}$  learn outcome  $y$

Regression: continuous outcome

Ex.  $y \in \mathbb{R}$   $\mathcal{L} = \sum_i (\hat{y}_i - y_i)^2$  squared error

Classification: discrete outcome

Ex.  $y \in \{0, 1\}$

- ▶ Logistic regression  $\hat{y} = \sigma(\mathbf{w}^T \mathbf{x})$
- ▶ SVM classifier

# Stochastic gradient descent

$\mathbf{w} \leftarrow \mathbf{w} - \gamma \nabla_{\mathbf{w}} \mathcal{L} \rightarrow$  gradient computed on a batch of data

```
from autograd import grad  
parameters -= GAMMA * grad(loss)(parameters)
```

## TP 2: Unsupervised learning

Learn  $\mathbf{x}$


- ▶  $K$ -nearest neighbors
- ▶ Dimensionality reduction (TP 3)
- ▶ Latent factor model (didn't have time)

# TP 3: Reinforcement learning

Dynamic programming's *Principle of Optimality* (Bellman, 1952)

value function  $v_{\pi}(s)$   
action-value function  $q_{\pi}(s, a)$  } for policy  $\pi : S \rightarrow A \rightarrow R \rightarrow S' \rightarrow A'$

Q-learning (1992), SARSA (1996), policy gradient (2000)



## TP 4: Deep learning

- ▶ Architecture of layers (CNN, dense) from input to output
- ▶ Ex. dense  $n \rightarrow d_1 \rightarrow d_2 \rightarrow \dots \rightarrow 1$

Deep neural networks are universal function approximators.

### History

- ▶ 1958: Perceptron (Rosenblatt)
- ▶ 2012: ImageNet recognition (Krizhevsky, Sutskever & Hinton)
- ▶ 2014: GANs image generation (Goodfellow et al.)
- ▶ 2016: AlphaGo beats Go world champion (DeepMind)
- ▶ 2016: WaveNet speech synthesis (DeepMind)
- ▶ 2017: AlphaGo Zero

# Project: Roomba

Choose features  $\mathbf{x}(S, A)$

Simplest (position), or **tile coding**

Choose function approximation

$Q(S, A) = \mathbf{w}^T \mathbf{x}(S, A)$  or *DeepNet*( $\mathbf{x}(S, A)$ )

Choose training algorithm

- ▶ Try a random policy, pick random  $\mathbf{w}$
- ▶ SARSA with tile coding  $\mathbf{x}(S, A)$
- ▶ Policy gradient?